

ST300

Project

Candidate Number: 37443

Session: 2021/22

1 Analysis

1.1 Data description

The dataset analysed is the "Countries of the World" obtained on *Kaggle*, containing 227 rows of data on economic and geographic variables of the countries and regions of the world, compiled between 2004 and 2010, as different countries perform census at different years, and certain types of data is only compiled every 5 or 10 years. The task is to propose model to explain the relationship between GDP per capita and the other variables in the dataset. We note that 48 rows contain missing values in one or more columns, which will be excluded in this analysis. This dataset contains the following variables:

Independent categorical variables: **Region**, which has 10 levels (originally 11, Baltic countries (of which there are 3) are recategorised as Eastern European countries) and **Climate** which has 6 levels. (Level 1 - Dry tropical or ice; 2 - Wet tropical; 3 - Temperate humid subtropical; 4 - Dry hot summers and wet winters. Levels in between (e.g. 1.5) also exist.)

Independent continuous variables: **Population**, **Area_sqmi**, **Popdensity_persqmi**, **Coastline**, **Netmigration**, **Infantmortality_per1k**, **Literacy**, **Phones_per1k**, **Crops**, **Other**, **Birthrate**, **Deathrate**, **Agriculture**, **Industry**, **Service**. Note: **Crops** and **Other** sum to 1, where **Crops** represent the proportion of land cultivated for crops (including permanent crops); similarly **Agriculture**, **Industry** and **Service** represent economic sector composition and sum to 1.

Dependent (continuous) variable: **GDP_pc**.

1.2 Exploratory data analysis




Figure 1: Histogram of GDP per capita

Firstly, from the left-hand histogram in Figure 1 we observe that the data is heavily right-skewed, which is to be expected in the context of GDP per capita of different countries. Hence, we perform a logarithmic transformation on the dependent variable to obtain a better spread of outcome data values. The results is shown in right-hand histogram in the same figure.

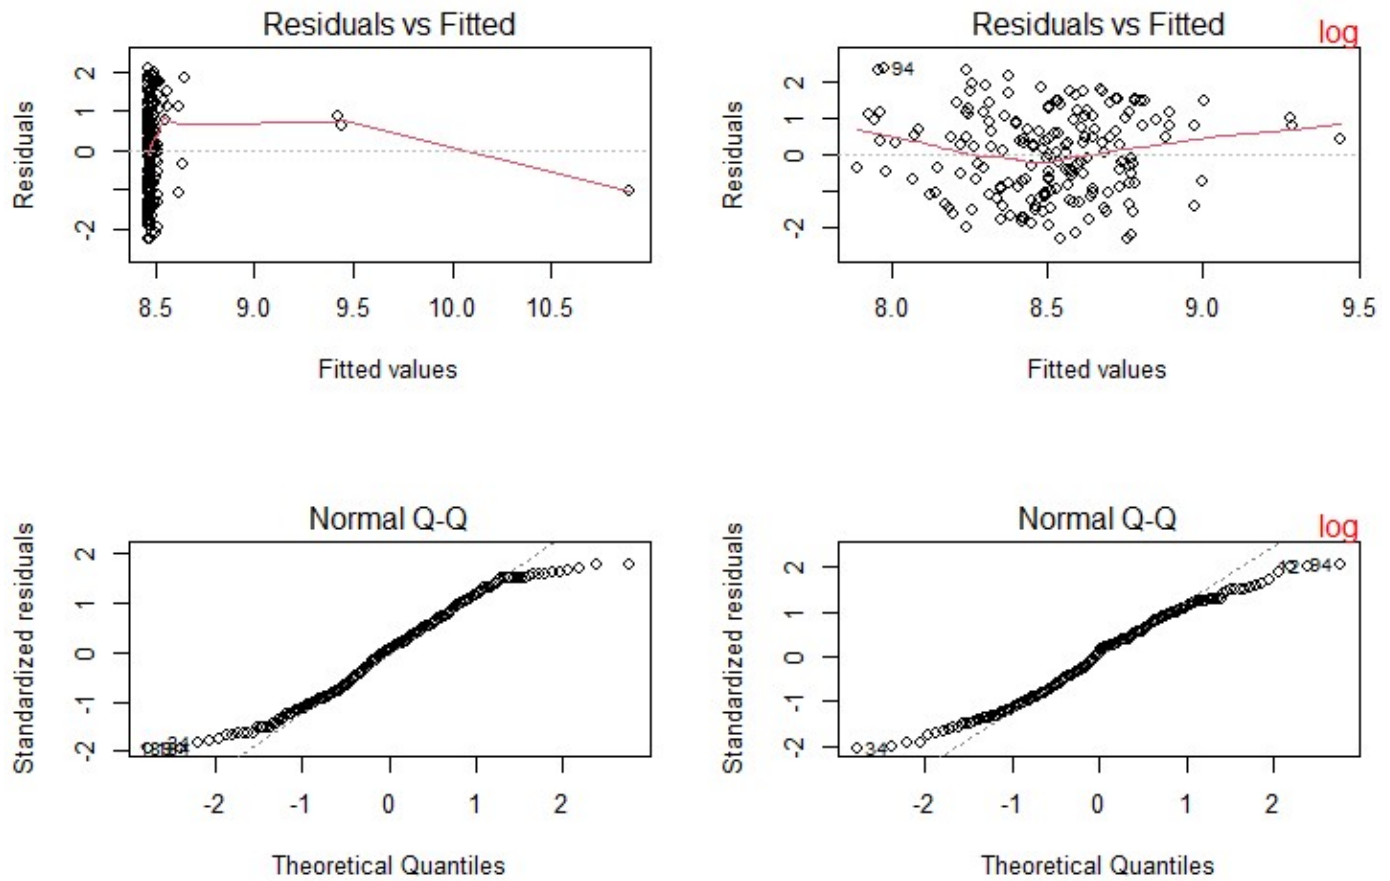


Figure 2: Comparison of fit under logarithmic transformation

We then identify the relationship between $\log(\text{GDP_pc})$ and each of the continuous variables separately. Table 1 in the appendix compares the correlation coefficient before and after a logarithmic transformation. In some cases, the correlation coefficient increases in magnitude significantly, for example for **population** it has increased in magnitude from -0.0141 to -0.1651 under a logarithmic transformation; whereas in other cases it has led to a decrease in magnitude, for example **Literacy** decreased in magnitude from 0.6848 to 0.6499.

Figure 2 shows examples of improved fit after logarithmic transformations, indicating significantly better homoscedasticity as shown by the residuals-fitted plot, and normality deviation remains at an acceptable level as verified by the normal Q-Q plot. The two plots on the left refer to **Popdensity_persqmi**, and the two plots on the right refer to $\log(\text{Popdensity_persqmi})$.

1.3 Variable transformation

Upon examining the relationship between $\log(\text{GDP_pc})$ and each of the continuous variables separately, subject to normality and homoscedasticity, as verified by the Residuals-Fitted and Normal Q-Q plots, **Population**, **Area_sqmi**, **Popdensity_persqmi**, **Infantmortality_per1k**, **Phones_per1k** will be transformed logarithmically.

As shown in Table 1, after the transformation, the correlation coefficient for these variables has increased in magnitude. Note that **Area_sqmi** has a positive but near zero r -value before the transformation (0.0465), and larger but a negative r -value after the transformation (-0.2543).

1.4 Variable selection

We initially attempt to fit the model with all of independent variables. Perhaps surprisingly, this model has adjusted R^2 value 0.8786, and a near zero p-value for the F-statistic. However, examining each of the coefficients, most have an insignificant p-value for the t-statistic, as well as large standard errors for the estimate.

To improve interpretability of this model, we attempt to improve the significance of these parameters, by only including a subset of the variables. We perform variable selection, using forward selection (FS) and backward elimination (BE) algorithms. The models produced by both algorithms are in agreement, with eight variables excluded: **Population**, **Area_sqmi**, **Popdensity_persqmi**, **Coastline**, **Climate**, **Literacy**, **Industry** and **Service**.

We first compare the model with and without the seven continuous variables mentioned above (we will investigate the variable **Climate** separately). The ANOVA test yields a p-value of 0.8867 for the F-statistic, along with a slight improvement to the adjusted R^2 value 0.8818. Therefore, we will drop these 7 variables. Next we investigate the categorical variable **Climate**. Again using an ANOVA test, the p-value is 0.7065 for the F-statistic, again with a slight improvement to the adjusted R^2 value 0.8833. Hence, we will also drop the variable **Climate**.

We note that many of the coefficients for the different levels of the categorical variable **Region** have relatively large standard errors and insignificant p-values for the t-statistic. An ANOVA test comparing models with and without the variable **Region** yields a p-value of 0.05037 for the F-statistic, which is weakly insignificant. However, we will retain this variable as otherwise we would have eliminated all categorical variables from our dataset.

Finally, we investigate potentially multicollinear variables. First, we observe that the variables **Crops** and **Other**, representing land use decomposition of a country. By definition these two variables sum to 1, and observing the correlation matrix in Appendix 3, these two variables are perfectly correlated (inversely), as expected. We compare the models with and without the variable **Other**. The ANOVA test yields a p-value of 0.01686 for the F-statistic. Hence we will drop the variable **Other**, as the value can be calculated from the other variable.

We also investigate possible multicollinearity between **Birthrate** and **Infantmortality_per1k**. Again from the correlation matrix in Appendix 3, the correlation coefficient between infant mortality rate and birthrate is 0.8389, suggesting strong relationship between these two variables. Intuitively, infant mortality rate relates to both deathrate and birthrate, as high infant mortality may lead to giving birth to more infants if fewer infants are expected to survive to adulthood, and similar contribute to deathrate if the population is relatively young with infant mortality contributing heavily to the population deathrate. This effect is expected to be more prominent in least-developed and developing countries, and less so in developed countries, where healthcare systems may be better. Again, we perform an ANOVA test comparing the models with and without both **Birthrate** and **Deathrate**. We obtain a p-value 0.0032 for the F-test, so we retain this variable.

We then also examine multicollinearity using variance inflation factor (VIF) calculated by the **cars** package in R. The full table is presented in Table 2 in the appendix. All GVIF values are below 10. Hence, there is no significant concern of multicollinearity in this model.

1.5 Model diagnostics

Finally, we check the model for potential outliers by calculating Cook's distance, as shown in Figure 3. The red line represents a threshold of $\frac{4}{n} \approx 0.02$. In this plot, there are 14 observations that exceed this threshold. However, re-fitting the model without these points does indeed increase the adjusted R^2 value. (0.9137 without vs 0.8798 including these observations).

Intuitively, outliers are to be expected due to the nature of the dataset, as the population size $n \approx 200$ is relative small. However, removing these observations may cause the model to be less representative since the dataset contained data of all countries.

For example, the p-value for the t-statistic for the variable **Crops** has increased significantly from 0.0112 to 0.1154. Also if we removing these 14 observations, in addition to the 48 observations already removed due to incomplete date, we would have removed just under 30% of the total number of observations, which is not ideal. Hence, we will retain all these observations to maintain the representativeness of the model.




Figure 3: Cook's distance plot

1.6 Checking model assumptions

Finally, we verify the model assumptions are indeed not violated. From the left plot of Figure 4, the standardised residuals appears to be scattered randomly above and below the zero line, and all lie within ± 2 , suggesting constant variance of residuals. From the centre plot, all observations lie nowhere near the Cook's distance line, verifying that the outliers that we identified in the previous section is not influential. From the right plot, although there is slightly bending tail on one end, it remains at an acceptable level so we confirm that normality is satisfied. Finally, we calculate the sum of residual errors to be $-3.3411 \times 10^{-15} \approx 0$, and mean of residual errors to be $-1.8677 \times 10^{-17} \approx 0$. Overall, we are convinced that the assumptions of a linear regression model are satisfied.




Figure 4: Residuals-fitted, residuals-leverage and normal Q-Q plots

1.7 Final model

The model proposed is as follows:

$$\log(\text{GDPpc}) = 9.2763 + \alpha + 0.0232 \times \text{Nmig} - 0.3591 \times \log(\text{Inmor}) + 0.2004 \times \log(\text{Phone}) + \\ - 0.5441 \times \text{Crops} - 0.0223 \times \text{Birth} + 0.0240 \times \text{Death} - 1.2582 \times \text{Agri}$$

where α is a sum of indicator functions for the categorical variable **Region**, given below.

Note: most variable names abbreviated. Nmig - **Netmigration**, Inmor - **Infantmortality_per1k**, Phone - **Phones_per1k**, Birth - **Birthrate**, Death - **Deathrate**, Agri - **Agriculture**.

$$\alpha = 0.0483 \times \mathbf{1}_{\text{Asia}} - 0.2766 \times \mathbf{1}_{\text{CIS}} - 0.0477 \times \mathbf{1}_{\text{E.Europe}} + 0.0332 \times \mathbf{1}_{\text{L.America}} + 0.2685 \times \mathbf{1}_{\text{NearEast}} + \\ + 0.1397 \times \mathbf{1}_{\text{N.Africa}} + 0.7462 \times \mathbf{1}_{\text{N.America}} + 0.1661 \times \mathbf{1}_{\text{Oceania}} + 0.3296 \times \mathbf{1}_{\text{W.Europe}}$$

Note: reference category Sub-Saharan Africa; Asia - Asia (Ex. Near East); CIS - Commonwealth of Independence States (mainly former Soviet Union states); L.America - Latin American and the Caribbean; E.Europe - Eastern Europe; N.Africa - Northern Africa; N.America - Northern America; W.Europe = Western Europe.

Observing the results shown in Appendix A.4, we note that this model has an adjust R^2 value of 0.8798, a F-statistic value of 82.44 with the associated p-value of less than $2.2 \times 10^{-16} \approx 0$. We also note that most of the variables are (or nearly) statistically significant (≤ 0.05), although for some of the categorical variables, some of the levels have relatively high p-values for the t-statistic. Nonetheless, the magnitude of standard errors of many of the coefficients is also much smaller (by one or more orders of magnitude) relative to the estimated coefficient. A summary is presented in Appendix A.4.

1.8 Interpretation of model parameters

Regarding the continuous variable **Netmigration**, defined as $\frac{\text{Number of Immigrants} - \text{Emigrants}}{0.5 \times (\text{Pop. at start of year} + \text{at end of year})} \times 1000$: keeping all other variables fixed, an 1 unit increase in the net migration rate, on average, leads to a change in GDP per capita by a factor of $e^{0.0232} = 1.0235$, equivalent to 2.35% increase.

Regarding the log-transformed continuous variable **Infantmortality_per1k**, defined as the deaths of children under one year of age per 1,000 live births: keeping all other variables fixed, a 1% increase in the infant mortality rate, on average, leads to a change in GDP per capita by a factor of $1.01^{-0.3591} = 0.9964$, equivalent to 0.36% decrease.

Regarding the log-transformed continuous variable **Phones_per1k**, defined as mobile phones per 1,000 people: keeping all other variables fixed, a 0.01 unit (or 1%) increase in the rate, on average, leads to a change in GDP per capita by a factor of $1.01^{0.2004} = 1.0020$, equivalent to 0.20% increase.

Regarding the continuous variable **Crops**, defined as the proportion of land cultivated for growing crops: keeping all other variables fixed, a 0.01 unit (or 1%) increase in proportion of land for permanent crops, on average, leads to a change in GDP per capita by a factor of $e^{0.01 \times -0.5441} = 0.9946$, equivalent to 0.54% decrease.

Regarding the continuous variable **Birthrate**, defined as total number of live births per 1,000 population: keeping all other variables fixed, a 0.01 unit (or 1%) increase in the birthrate, on average, leads to a change in GDP per capita by a factor of $e^{0.01 \times -0.0223} = 0.9998$, equivalent to 0.02% decrease.

Regarding the continuous variable **Deathrate**, defined as the number of deaths per 1,000 population: keeping all other variables fixed, a 0.01 unit (or 1%) increase in the deathrate, on average, leads to a change in GDP per capita by a factor of $e^{0.01 \times 0.0240} = 1.0002$, equivalent to 0.02% increase.

Regarding the continuous variable **Agriculture**, defined as the proportion of GDP per capita generated by the Agricultural (or primary) sector: keeping all other variables fixed, a 0.01 unit (or 1%) increase in the proportion, on average, leads to a change in GDP per capita by a factor of $e^{0.01 \times -1.2582} = 0.9875$, equivalent to 1.25% decrease.

Regarding the categorical variable **Region**, compared to a country being in Sub-Saharan Africa, keeping all other variables fixed, a country in:

- Asia, on average, has GDP per capita higher by 4.95% ($e^{0.0483} = 1.0495$);
- the CIS, on average, has GDP per capita lower by 24.16% ($e^{-0.2766} = 0.7584$);
- Eastern Europe, on average, has GDP per capita lower by 4.66% ($e^{-0.0477} = 0.9534$);
- Latin America & Caribbean, on average, has GDP per capita higher by 3.38% ($e^{0.0332} = 1.0338$);
- Near East, on average, has GDP per capita higher by 30.80% ($e^{0.2685} = 1.3080$);
- Northern Africa, on average, has GDP per capita higher by 14.99% ($e^{0.1397} = 1.1499$);
- North America, on average, has GDP per capita higher by 110.90% ($e^{0.7462} = 2.1090$);
- Oceania, on average, has GDP per capita higher by 18.07% ($e^{0.1661} = 1.1807$);
- Western Europe, on average, has GDP per capita higher by 39.04% ($e^{0.3296} = 1.3904$).

Judging by the sign of the coefficients, most seem to be in agreement with economic theory. For example, a higher net migration rate, is indicative of a country being attractive to people from other countries, which suggests that the country is more economically developed. Another example is that a lower infant mortality rate is indicative of better healthcare system, a common feature in more economically developed countries.

Finally, we observe that the variable **Agriculture** to be the most important continuous variable in this model. Again, according to the three-sector model in economic theory, economies generally shift its main focus from primary (agriculture) to secondary (industry), and finally to the tertiary sector (service). Thus, an agricultural based economy is strongly indicative of lower GDP per capita, again supporting the fact that this variable has the largest coefficient in magnitude.

1.9 Model limitations and conclusion

The main limitation of this model is the size being relatively small ≈ 200 , meaning that each individual observations may have a high influence on the estimated coefficients. In section 1.5 where we evaluated the Cook's distance plot, we decided to keep all the observations despite high Cook's distance, since we prefer the model being more representative. However, since changes in the variables measured in this dataset may change significantly over the years, the model will very likely be significantly different when comparing economic data from different years. Furthermore, the assumption that observations being independent may not be entirely valid, since the economy of one country may be significantly dependent on other countries, for example oil-producing countries.

However, many of the variables were excluded in the final model. Although a simpler model improves the explainability of the model, some of the variables which intuitively may have some relationship to GDP per capita according to economic theory is excluded, for example **Literacy**, where economic theory suggests that a higher literacy rate is indicative of higher productivity, hence higher GDP per capita. Nonetheless, the interpretation of the model given above also mostly agree with economic theory. Overall, we conclude that the model is of a good fit of the data, and also can be interpreted easily to draw relationships regarding GDP per capita and the other variables in the dataset.

A Appendix

A.1 Comparision of correlation coefficients under log-transformation

This table shows the correlation coefficient between the dependent variables $\log(\text{GDP_pc})$ and each of the continuous independent variables. The correlation matrix between each of the continuous independent variables is given in Appendix 3.

Where the table cell is empty, this means the transformation is inappropriate for this variable, due to zero values being sent to $-\infty$. Note, we could have also used a shifted logarithmic transformation, e.g. $\log(x + \epsilon)$. However, this is undesirable as a number of these variables have values in $[0, 1]$; and also that for very small ϵ , the magnitude of $\log(\epsilon)$ becomes increasingly large and negative; and futher, makes the model more complex to interpret.

	Raw	Log-transform
Population	-0.0141	-0.1651
Area_sqmi	0.0465	-0.2543
Popdensity_persqmi	0.1759	0.2165
Coastline	0.0466	
Netmigration	0.2396	
Infantmortality_per1k	-0.8292	-0.8838
Literacy	0.6848	0.6499
Phones_per1k	0.8473	0.8682
Crops	-0.0317	
Other	0.0317	0.0522
Birthrate	-0.8341	-0.8324
Deathrate	-0.3968	-0.4006
Agriculture	-0.7852	
Industry	0.1477	0.1339
Service	0.5913	0.5250

Table 1: Comparision of correlation coefficients under log-transformation

A.2 Table of generalised variance inflation factor (GVIF)

Table only include variables used in the final model.

	GVIF	Df	GVIF ^{1/(2*Df)}
Region	11.262875	9	1.143997
Netmigration	1.335901	1	1.155812
Infantmortality_per1k	7.100051	1	2.664592
Phones_per1k	8.133274	1	2.851890
Crops	1.240374	1	1.113721
Birthrate	7.544744	1	2.746770
Deathrate	2.621064	1	1.618970
Agriculture	3.009943	1	1.734919

Table 2: Table of generalised variance inflation factor (GVIF)

A.3 Correlation matrix of independent continuous variables

Correlation calculated after applying log-transformation to **Population**, **Area_sqmi**, **Popdensity_persqmi**, **Infantmortality_per1k** and **Phones_per1k**.

	Pop	Area	Popdn	Coast	Nmig	Inmor	Lit	Phone	Crops	Other	Birth	Death	Agri	Ind	Ser
Pop	1.0000	0.8359	-0.0296	-0.3416	0.0779	0.1666	-0.1884	-0.2569	0.0588	-0.0588	0.1062	0.1520	0.1657	0.2066	-0.3220
Area	0.8359	1.0000	-0.5734	-0.3875	0.0358	0.3078	-0.2340	-0.3530	-0.2104	0.2104	0.2429	0.2634	0.2532	0.2372	-0.4304
Popdn	-0.0296	-0.5734	1.0000	0.1957	0.0510	-0.3117	0.1447	0.2593	0.4710	-0.4710	-0.2837	-0.2529	-0.2135	-0.1236	0.3031
Coast	-0.3416	-0.3875	0.1957	1.0000	-0.2416	-0.0711	0.0996	0.1248	0.1371	-0.1371	-0.0635	-0.1486	-0.0323	-0.1890	0.1900
Nmig	0.0779	0.0358	0.0510	-0.2416	1.0000	-0.1643	-0.0538	0.0457	-0.2574	0.2574	-0.0351	0.0428	-0.0966	-0.0044	0.0915
Inmor	0.1666	0.3078	-0.3117	-0.0711	-0.1643	1.0000	-0.7036	-0.8492	-0.0902	0.0902	0.8389	0.4991	0.7184	-0.0151	-0.6428
Lit	-0.1884	-0.2340	0.1447	0.0996	-0.0538	-0.7036	1.0000	0.7538	0.1012	-0.1012	-0.7883	-0.4017	-0.6205	0.1057	0.4744
Phone	-0.2569	-0.3530	0.2593	0.1248	0.0457	-0.8492	0.7538	1.0000	0.0793	-0.0793	-0.8803	-0.5356	-0.7903	0.0937	0.6413
Crops	0.0588	-0.2104	0.4710	0.1371	-0.2574	-0.0902	0.1012	0.0793	1.0000	-1.0000	-0.1240	-0.0660	0.0271	-0.1223	0.0815
Other	-0.0588	0.2104	-0.4710	-0.1371	0.2574	0.0902	-0.1012	-0.0793	-1.0000	1.0000	0.1239	0.0660	-0.0271	0.1223	-0.0815
Birth	0.1062	0.2429	-0.2837	-0.0635	-0.0351	0.8389	-0.7883	-0.8803	-0.1240	0.1239	1.0000	0.4462	0.7040	-0.1205	-0.5417
Death	0.1520	0.2634	-0.2529	-0.1486	0.0428	0.4991	-0.4017	-0.5356	-0.0660	0.0660	0.4462	1.0000	0.4164	-0.0126	-0.3662
Agri	0.1657	0.2532	-0.2135	-0.0323	-0.0966	0.7184	-0.6205	-0.7903	0.0271	-0.0271	0.7040	0.4164	1.0000	-0.3528	-0.6135
Ind	0.2066	0.2372	-0.1236	-0.1890	-0.0044	-0.0151	0.1057	0.0937	-0.1223	0.1223	-0.1205	-0.0126	-0.3528	1.0000	-0.5214
Ser	-0.3220	-0.4304	0.3031	0.1900	0.0915	-0.6428	0.4744	0.6413	0.0815	-0.0815	-0.5417	-0.3662	-0.6135	-0.5214	1.0000

Table 3: Correlation matrix of independent continuous variables

Note: most variable names abbreviated. Pop - **Population**, Area - **Area_sqmi**, Popdn - **Popdensity_persqmi**, Coast - **Coastline**, Nmig - **Netmigration**, Inmor - **Infantmortality_per1k**, Lit - **Literacy**, Phone - **Phones_per1k**, Birth - **Birthrate**, Death - **Deathrate**, Agri - **Agriculture**, Ind - **Industry**, Ser - **Service**.

A.4 Summary of final model

Call:

```
lm(formula = log(GDP_pc) ~ ., data = df3)
```

Residuals:

Min	1Q	Median	3Q	Max
-1.16391	-0.25659	0.03987	0.25698	1.05491

Coefficients:

	Estimate	Std. Error	t value	Pr(> t)	
(Intercept)	9.276271	0.447938	20.709	< 2e-16	***
RegionASIA (EX. NEAR EAST)	0.048340	0.137986	0.350	0.726549	
RegionC.W. OF IND. STATES	-0.276615	0.176866	-1.564	0.119775	
RegionEASTERN EUROPE	-0.047688	0.196729	-0.242	0.808772	
RegionLATIN AMER. & CARIB	0.033156	0.134579	0.246	0.805710	
RegionNEAR EAST	0.268535	0.174399	1.540	0.125566	
RegionNORTHERN AFRICA	0.139740	0.273489	0.511	0.610081	
RegionNORTHERN AMERICA	0.746222	0.322059	2.317	0.021754	*
RegionOCEANIA	0.166085	0.164900	1.007	0.315349	
RegionWESTERN EUROPE	0.329576	0.175376	1.879	0.062007	.
Netmigration	0.023170	0.007461	3.105	0.002244	**
Infantmortality_per1k	-0.359130	0.076474	-4.696	5.62e-06	***
Phones_per1k	0.200418	0.051113	3.921	0.000130	***
Crops	-0.544097	0.212107	-2.565	0.011219	*
Birthrate	-0.022266	0.007464	-2.983	0.003294	**
Deathrate	0.024035	0.009531	2.522	0.012637	*
Agriculture	-1.258186	0.351594	-3.579	0.000456	***

Signif. codes: 0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

Residual standard error: 0.4092 on 162 degrees of freedom

Multiple R-squared: 0.8906, Adjusted R-squared: 0.8798

F-statistic: 82.44 on 16 and 162 DF, p-value: < 2.2e-16

B R code

```

1 #Setup
2 library(ggplot2)
3 library(xtable)
4 library(dplyr)
5 library(stringr)
6 library(data.table)
7 df0 <- read.csv("countries of the world.csv", dec=",", strip.white=T, head(
  T))
8
9 #Data clean-up
10 dim(df0[rowSums(is.na(df0))>0,]) #Check for rows containing missing values
11 #48 out of 227 rows have missing values in one or more columns
12 df <- na.exclude(df0) #we will remove these rows
13
14 df$Region <- str_trim(df$Region) #remove trailing white spaces
15
16 #rename columns
17 oldnames <- c("Area..sq..mi..", "Pop..Density..per.sq..mi..", "Coastline..
  coast.area.ratio.", "Net.migration", "Infant.mortality..per.1000.births.",
  "GDP....per.capita.", "Literacy....", "Phones..per.1000.", "Arable....", "
  Crops....", "Other....")
18 newnames <- c("Area_sqmi", "Popdensity_persqmi", "Coastline", "Netmigration", "
  Infantmortality_per1k", "GDP_pc", "Literacy", "Phones_per1k", "Arableland", "
  Crops", "Other")
19 df <- setnames(df, old=oldnames, new=newnames, skip_absent=T)
20
21 #rescale literacy and land usage composition from percentage to decimal:
22 index <- c(10, 12:14)
23 df[, index] <- df[, index]/100
24
25 #check factor with highest number of observations
26 table(df$Region)
27 table(df$Climate)
28
29 #Categories of low counts (after observations with missing data):
30 #N.America-2: Bermuda, USA
31 #Baltics-2: Estonia, Latvia
32 #N.Africa-3: Algeria, Egypt, Tunisia
33 #N.America more economically different to rest of America, retain
34 #N.Africa more historically and culturally different to Sub-Saharan Africa,
  retain
35 #However will merge Baltics into Eastern Europe
36 df$Region[df$Region == "BALTICS"] <- "EASTERN EUROPE"
37
38 #Convert categorical variables to factors
39 df$Region <- as.factor(df$Region)
40 df$Climate <- as.factor(df$Climate)
41
42 #Redefine reference level to one with most observations:
43 df$Climate <- relevel(df$Climate, "2")
44 df$Region <- relevel(df$Region, "SUB-SAHARAN AFRICA")
45
46 #Combine Arableland and Crops
47 #Arablelands = Land cultivated for crops

```

```

48 #Crops = Permenant Crops
49 #Other = Other land use not arableland nor crops
50 df$Crops <- df$Arableland + df$Crops #combine crops
51 df <- subset(df,select=-c(Arableland)) #remove
52
53 #Histogram of GDP_pc
54 par(mfrow=c(1,2))
55 hist(df$GDP_pc, xlab="GDP per capita", main="Histogram of GDP per capita")
56 hist(log(df$GDP_pc), xlab="log(GDP per capita)", main="Histogram of log(GDP
    per capita)")
57
58 #index for continuous variables
59 variable_index <- c(3:8,10:13,15:19)
60
61 #Exploring the relationship between GDP_pc and each continuous independent
    variable separately
62 for (i in variable_index){
63   plot1 <- ggplot(aes_string(x=names(df)[i],y=log(df$GDP_pc)),data=df)+
64     geom_point()+
65     ylab("log(GDPperCap)")+
66     geom_smooth(method="lm",se=F)
67   print(plot1)
68 }
69
70 #Correlation between GDP_pc and continuous independent variables, comparing
    with and without log-transform
71 cor0 <- t(cor(log(df$GDP_pc),df[,variable_index]))
72 cor1 <- t(cor(log(df$GDP_pc),log(df[,variable_index])))
73 corr <- cbind(cor0,cor1)
74 colnames(corr) <- c("Raw", "Log-transform")
75 xtable(corr, digits=4) #table output
76
77 #index for continuous variables that can be log-transformed
78 logtransformable <- setdiff(variable_index,c(6:7,10,13:14,17))
79
80 #Standardised Residuals and Normal Q-Q plot
81 for (i in variable_index){
82   par(mfrow=c(2,2),oma=c(0,0,2,0)) #top wide margin
83   plot.new()
84   mtext(names(df)[i],outer=TRUE,cex=1.5) #plot title
85   model_1 <- lm(paste("log(GDP_pc)~",names(df)[i]),data=df)
86   par(mfg=c(1,1)) #plot at top-left
87   plot(model_1,1,sub.caption = "") #Standardised residuals
88   par(mfg=c(2,1)) #plot at bottom-left
89   plot(model_1,2,sub.caption = "") #Normal Q-Q
90   #compare with graphs if log-transforming the variable
91   if (is.element(i,logtransformable)){
92     model_2 <- lm(paste("log(GDP_pc)~log(",names(df)[i],")"),data=df)
93     par(mfg=c(1,2)) #plot at top-right
94     plot(model_2,1,sub.caption = "") #Normal Q-Q
95     mtext("log",adj=1, col="red") #plot label
96     par(mfg=c(2,2)) #plot at bottom-right
97     plot(model_2,2,sub.caption = "") #Normal Q-Q
98     mtext("log",adj=1, col="red") #plot label
99   }
100 }

```

```

101
102 #Transforming variables (Log-transform)
103 logvar <- c(3:5,8,11)
104 df[logvar] <- log(df[logvar])
105
106 #Correlation matrix of continuous independent variables
107 cor2 <- cor(df[,variable_index])
108 xtable(cor2,digits=4) #table output
109
110 #Verifying relationship after transformation
111 for (i in variable_index){
112   plot2 <- ggplot(aes_string(x=names(df)[i],y=df$GDP_pc),data=df)+
113     geom_point()+
114     ylab("log(GDPperCap)") +
115     geom_smooth(method="lm",se=F)
116   print(plot2)
117 }
118
119 #Initial fit with all variables (excluding Country as this is a label)
120 Full<- lm(log(GDP_pc) ~.-Country,data=df)
121 summary(Full) #adjusted r2 = 0.8786
122
123 Null <- lm(log(GDP_pc) ~1,data=df) #Model with only the intercept
124
125 #Backward elimination algorithm:
126 BE <- step(Full,scope=list(lower=Null,upper=Full),direction="backward",
127   trace=F)
128 summary(BE)
129 #Coefficients dropped: Population, Area_sqmi, Popdensity_persqmi, Coastline
130   , Literacy, Climate, Industry, Service
131
132 #Forward selection algorithm:
133 FS <- step(Null,scope=list(lower=Null,upper=Full),direction="forward",trace
134   =F)
135 summary(FS)
136 #Coefficients dropped: Population, Area_sqmi, Popdesnity_persqmi, Coastline
137   , Literacy, Climate, Industry, Service
138
139 #First revision: Drop continuous variables not included in both BE and FS
140 #Drop Population, Area_sqmi, Popdensity_persqmi, Coastline, Literacy,
141   Industry, Service
142 df1 <- subset(df,select=-c(Country,Population,Area_sqmi,Popdensity_persqmi,
143   Coastline,Literacy,Industry,Service))
144 firstmodel <- lm(log(GDP_pc) ~ .,data=df1)
145 summary(firstmodel) #adjusted r2 = 0.8818
146 anova(firstmodel,Full) #p-value 0.8867 > 0.05
147
148 #Test dropping Climate:
149 secondmodel <- lm(log(GDP_pc) ~ .-Climate,data=df1)
150 summary(secondmodel) #adjusted r2 = 0.8833
151 anova(secondmodel,firstmodel) #p-value 0.7065 > 0.05
152 df2 <- subset(df1,select=-c(Climate)) #remove
153
154 #Test dropping Region:
155 thirdmodel <- lm(log(GDP_pc) ~ .-Region,data=df2)
156 summary(thirdmodel) #adjusted r2 = 0.8775

```

```
151 anova(thirdmodel,secondmodel) #p-value 0.05037 ~ 0.05 #keep
152
153 #Test dropping Other:
154 fourthmodel <- lm(log(GDP_pc) ~ .-Other,data=df2)
155 summary(fourthmodel) #adjusted r2 = 0.8798
156 anova(fourthmodel,secondmodel) #p-value 0.01686 < 0.05
157 df3 <- subset(df2,select=-c(Other)) #remove
158
159 #Test dropping Birthrate:
160 fifthmodel <- lm(log(GDP_pc) ~ .-Birthrate,data=df3)
161 summary(fifthmodel) #adjusted r2 = 0.874
162 anova(fifthmodel,fourthmodel) #p-value 0.0032 < 0.05 #keep
163
164 #final model
165 finalmodel <- lm(log(GDP_pc) ~ .,data=df3)
166 summary(finalmodel)
167
168 #Check for multi-collinearity
169 car::vif(finalmodel)
170 xtable(car::vif(finalmodel),digits=c(0,6,0,6))
171 #all variables have VIF < 10, so no indication of multi-collinearity
172
173 par(mfrow=c(1,3)) #graph output
174 plot(finalmodel,1,sub.caption = "") #Standardised residuals plot
175 plot(finalmodel,5,sub.caption = "") #Residuals-Leverage plot
176 plot(finalmodel,2,sub.caption = "") #Normal Q-Q plot
177
178 par(mfrow=c(1,1)) #graph output
179 plot(finalmodel,4,sub.caption = "") #Cook's distance plot
180 abline(h=0.02, col="red", lty=2) #threshold 4/n approx 0.02
181
182 #Compute new model without outliers:
183 cd <- cooks.distance(finalmodel)
184 outliers <- names(cd)[cd > 0.02] #15 outliers
185 df4 <- df3[!(row.names(df) %in% outliers),]
186 finalmodel2 <- lm(log(GDP_pc) ~ .,data=df4)
187 summary(finalmodel2)
188
189 mean(finalmodel$residuals) #mean of residuals: near zero
190 sum(finalmodel$residuals) #sum of residuals: near zero
```